

SIMATS ENGINEERING

CO3-ASSESSMENT TOOL2-MCQ

COURSE NAME: Artificial Intelligence

COURSE CODE: CSA1709

COURSE OUTCOME COVERED

CO3: Analyze knowledge representation and reasoning using propositional logic, FOL, inference, resolution, chaining methods, knowledge representation to perform logical reasoning for drawing valid conclusions and decision making. (BL4)

Assessment Tool Weightage: 10%

QUESTIONS: 30

Total Marks:30

Code of Conduct:

I M.Indupriya(192372085)) certify that this submission is my original work and that I have adhered to the guidelines specified for this assessment.

I understand that any violation of academic integrity rules will result in disciplinary action.

indupriya

Signature of the student

Circle the most appropriate option for each question.

Q1. Which of the following best describes a Logical Agent in AI?

- a) An agent that reacts randomly to the environment
- b) An agent that uses a knowledge base and logical inference to decide actions
- c) An agent that only follows pre-defined scripts
- d) An agent that learns from rewards

Q2. In Propositional Logic, ' $P \rightarrow Q$ ' is logically equivalent to:

- a) $P \wedge \neg Q$
- b) $\neg P \vee Q$
- c) $\neg Q \rightarrow \neg P$ only
- d) Both b and c

Q3. Conjunctive Normal Form (CNF) requires the formula to be expressed as:

- a) A disjunction of conjunctions
- b) A conjunction of disjunctions (clauses)
- c) A conjunction of literals only
- d) A sequence of implications

Q4. In FOL, the sentence $\forall x: \text{Student}(x) \rightarrow \text{Passes}(x)$ means:

- a) Some students pass
- b) Every entity that is a student passes
- c) If something passes, it is a student
- d) There exists at least one student who passes

Q5. Unification in FOL is the process of finding substitutions that:

- a) Convert clauses to CNF
- b) Make two logical expressions syntactically identical
- c) Prove a goal using backward chaining
- d) Derive facts using Modus Ponens

Q6. In the Resolution algorithm, deriving the empty clause indicates:

- a) The knowledge base is consistent
- b) The negated goal is satisfiable
- c) A contradiction — the original goal is proven
- d) More clauses need to be added

Q7. Forward Chaining is best described as a strategy.

- a) Goal-driven / top-down
- b) Data-driven / bottom-up
- c) Heuristic-based depth-first
- d) Probabilistic inference

Q8. Backward Chaining begins reasoning from:

- a) All known facts toward a goal
b) The goal and works backward to verify supporting facts
c) A random clause in the knowledge base
d) The resolution

Q9. Which step is NOT part of converting a sentence to CNF?

- a) Eliminate biconditionals ($\leftrightarrow$) quantifiers
- b) Move negations inward using De Morgan's laws
- c) Apply Skolemization to remove existential quantifiers
- d) Distribute AND over OR

Q10. Knowledge Engineering in FOL involves:

- a) Training a neural network on domain-specific data c) Writing heuristic search strategies
b) Defining domain concepts, relations, and axioms d) Encoding reward functions for reinforcement learning formally in FOL

Q11. Modus Ponens states: If $P \rightarrow Q$ is known and P is true, then:

- a) $\neg Q$ is true c) P is uncertain
b) Q must be true d) The KB is inconsistent

Q12. The Most General Unifier (MGU) is the substitution that:

- a) Is the most specific possible match c) Eliminates all variables from a clause
b) Makes two terms identical with minimal variable d) Applies resolution to produce
CNF binding

Q13. Which inference method is more appropriate when a specific goal is known?

- a) Forward Chaining c) Backward Chaining
b) Truth Table method d) A* search

Q14. Propositional Logic differs from First Order Logic because it:

- a) Supports probabilistic reasoning c) Is always more expressive
b) Cannot represent objects, relations, or quantifiers — d) Handles uncertainty better only truth values of atomic propositions

Q15. Resolution Refutation proves a goal G by:

- a) Directly deriving G from the KB using forward chaining
- c) Checking G in the truth table
- b) Adding $\neg G$ to the KB and deriving a contradiction
- d) Applying Modus Tollens repeatedly (empty clause)

Q16. Learning from observation in AI refers to:

- a) An agent receiving numerical rewards from the symbolic environment
- c) Acquiring knowledge by following explicit rules
- b) Improving performance by examining examples or generalization experience from the environment
- d) Memorizing a fixed dataset without

Q17. Inductive learning is best described as:

- a) Deriving specific conclusions from general rules
- c) Memorizing training data exactly
- b) Generalizing a hypothesis from specific training policies examples
- d) Applying search algorithms to find optimal

Q18. In a Decision Tree, what does each internal node represent?

- a) A class label / output decision
- c) A test on an attribute / feature
- b) A probability distribution over outcomes
- d) A training data sample

Q19. Explanation-Based Learning (EBL) differs from inductive learning because it:

- a) Requires large amounts of training data to generalize and generalize
- c) Uses prior domain knowledge to explain a single example
- b) Learns only from numeric data using regression
- d) Applies unsupervised clustering to find hidden patterns

Q20. In statistical learning, a Naive Bayes classifier assumes:

- a) All features are conditionally dependent on each other independent given the given the class
- c) All features are conditionally class label
- b) The data must follow a uniform distribution
- d) Labels are determined by a reward function

Q21. Learning with complete data (no hidden variables) in a Bayesian network involves:

a) Using the EM algorithm to estimate missing values c) Directly computing maximum likelihood estimates from

observed data

b) Training a deep neural network with dropout weights regularization

d) Applying reinforcement signals to update

Q22. The Expectation–Maximization (EM) algorithm is used when:

a) Training data is fully labeled and no variables are c) The dataset contains hidden (latent) variables that hidden cannot be directly observed

b) An agent explores an unknown environment using trial d) Backpropagation fails to converge in a neural network and error

Q23. In a Neural Network, the activation function serves to:

a) Initialize the weights of each neuron to zero complex

c) Introduce non–linearity to enable learning of patterns

b) Normalize the input dataset before training

d) Store previously learned examples in memory

Q24. Which of the following correctly describes backpropagation in a neural network?

a) Weights are updated in the forward pass using input from output to input gradients

c) The error is propagated backward layers to adjust weights

b) Outputs of each layer are stored and reused in later each hidden layer epochs

d) A separate model is trained for independently

Q25. A Support Vector Machine (SVM) — a type of kernel machine — classifies data by:

a) Minimizing the number of training examples used c) Finding the maximum–margin hyperplane that separates classes

b) Grouping similar data points into clusters using environment centroids d) Updating weights using Q–values from an

Q26. The kernel trick in kernel machines (SVMs) allows:

a) Reducing the number of support vectors needed c) Operating in a high–dimensional feature space without

explicitly computing coordinates

b) Accelerating gradient descent convergence

d) Replacing hidden layers in a deep neural network

Q27. In Reinforcement Learning, the agent aims to maximize:

- a) Classification accuracy on a fixed test set
- b) The number of features selected from the training data
- c) Cumulative reward received over time through interaction with the environment
- d) The size of the neural network hidden layer

Q28. The Q-Learning algorithm in Reinforcement Learning is classified as:

- a) A supervised learning algorithm requiring labeled learning datasets
- b) An unsupervised clustering technique using centroids
- c) A model-free, off-policy temporal difference method
- d) A decision-tree-based ensemble classifier

Q29. Overfitting in machine learning occurs when:

- a) The model performs poorly on both training and test data
- b) The learning rate is set too low during gradient descent
- c) The model fits training data too closely and generalizes poorly to new data
- d) The training set has more features than examples

Q30. Which of the following best distinguishes Reinforcement Learning from Supervised Learning?

- a) RL uses labeled input-output pairs; SL uses reward feedback; SL signals
- b) Both RL and SL require explicit reward functions
- c) RL learns via trial-and-error with reward feedback; SL learns from labeled examples
- d) RL is always faster to converge than SL on any dataset

1. B — An agent that uses a knowledge base and logical inference to decide actions
2. D — Both b and c
3. B — A conjunction of disjunctions (clauses)
4. B — Every entity that is a student passes
5. B — Make two logical expressions syntactically identical
6. C — A contradiction — the original goal is proven
7. B — Data-driven / bottom-up
8. B — The goal works backward to verify supporting facts
9. C — Apply Skolemization to remove existential quantifiers
10. B — Defining domain concepts, relations, and axioms
11. B — Q must be true
12. B — Makes two terms identical with minimal variable binding
13. C — Backward Chaining
14. B — Cannot represent objects, relations, or quantifiers — only truth values of atomic propositions
15. B — Adding $\neg G$ to the KB and deriving a contradiction (empty clause)
16. B — Improving performance by examining examples or experience from the environment
17. B — Generalizing a hypothesis from specific training examples
18. C — A test on an attribute / feature
19. C — Uses prior domain knowledge to explain and generalize a single example
20. C — All features are conditionally independent given the class label
21. C — Directly computing maximum likelihood estimates from observed data
22. C — The dataset contains hidden (latent) variables that cannot be directly observed
23. C — Introduce non-linearity to enable learning of complex patterns
24. C — The error is propagated backward from output to input layers to adjust weights
25. C — Finding the maximum-margin hyperplane that separates classes
26. C — Operating in a high-dimensional feature space without explicitly computing coordinates
27. C — Cumulative reward received over time through interaction with the environment
28. C — A model-free, off-policy temporal difference learning method
29. C — The model fits training data too closely and generalizes poorly to new data
30. C — RL learns via trial-and-error with reward feedback; SL learns from labeled example

RUBRICS FOR CALCULATION

Criteria	Weightage (%)	Level 4 (Exemplary)	Level 3 (Proficient)	Level 2 (Developing)	Level 1 (Beginning)
Conceptual Understanding	30%	Demonstrates thorough understanding; answers all questions accurately	Shows good understanding with few errors	Partial understanding; several incorrect responses	Lacks understanding; majority of answers incorrect
Accuracy of Answers	25%	All answers are correct	Most answers are correct with minor mistakes	Some correct answers; frequent errors	Mostly incorrect answers
Application of Knowledge	20%	Correctly applies concepts to problem-based questions	Applies concepts with minor errors	Limited ability to apply concepts	Unable to apply concepts

Time Management	15%	Completes quiz well within time efficiently	Completes on time with minor delays	Requires extra time or rushes through questions	Unable to complete within given time
Question Interpretation	10%	Interprets all questions correctly and responds appropriately	Misinterprets a few questions	Often misinterprets questions	Unable to understand questions